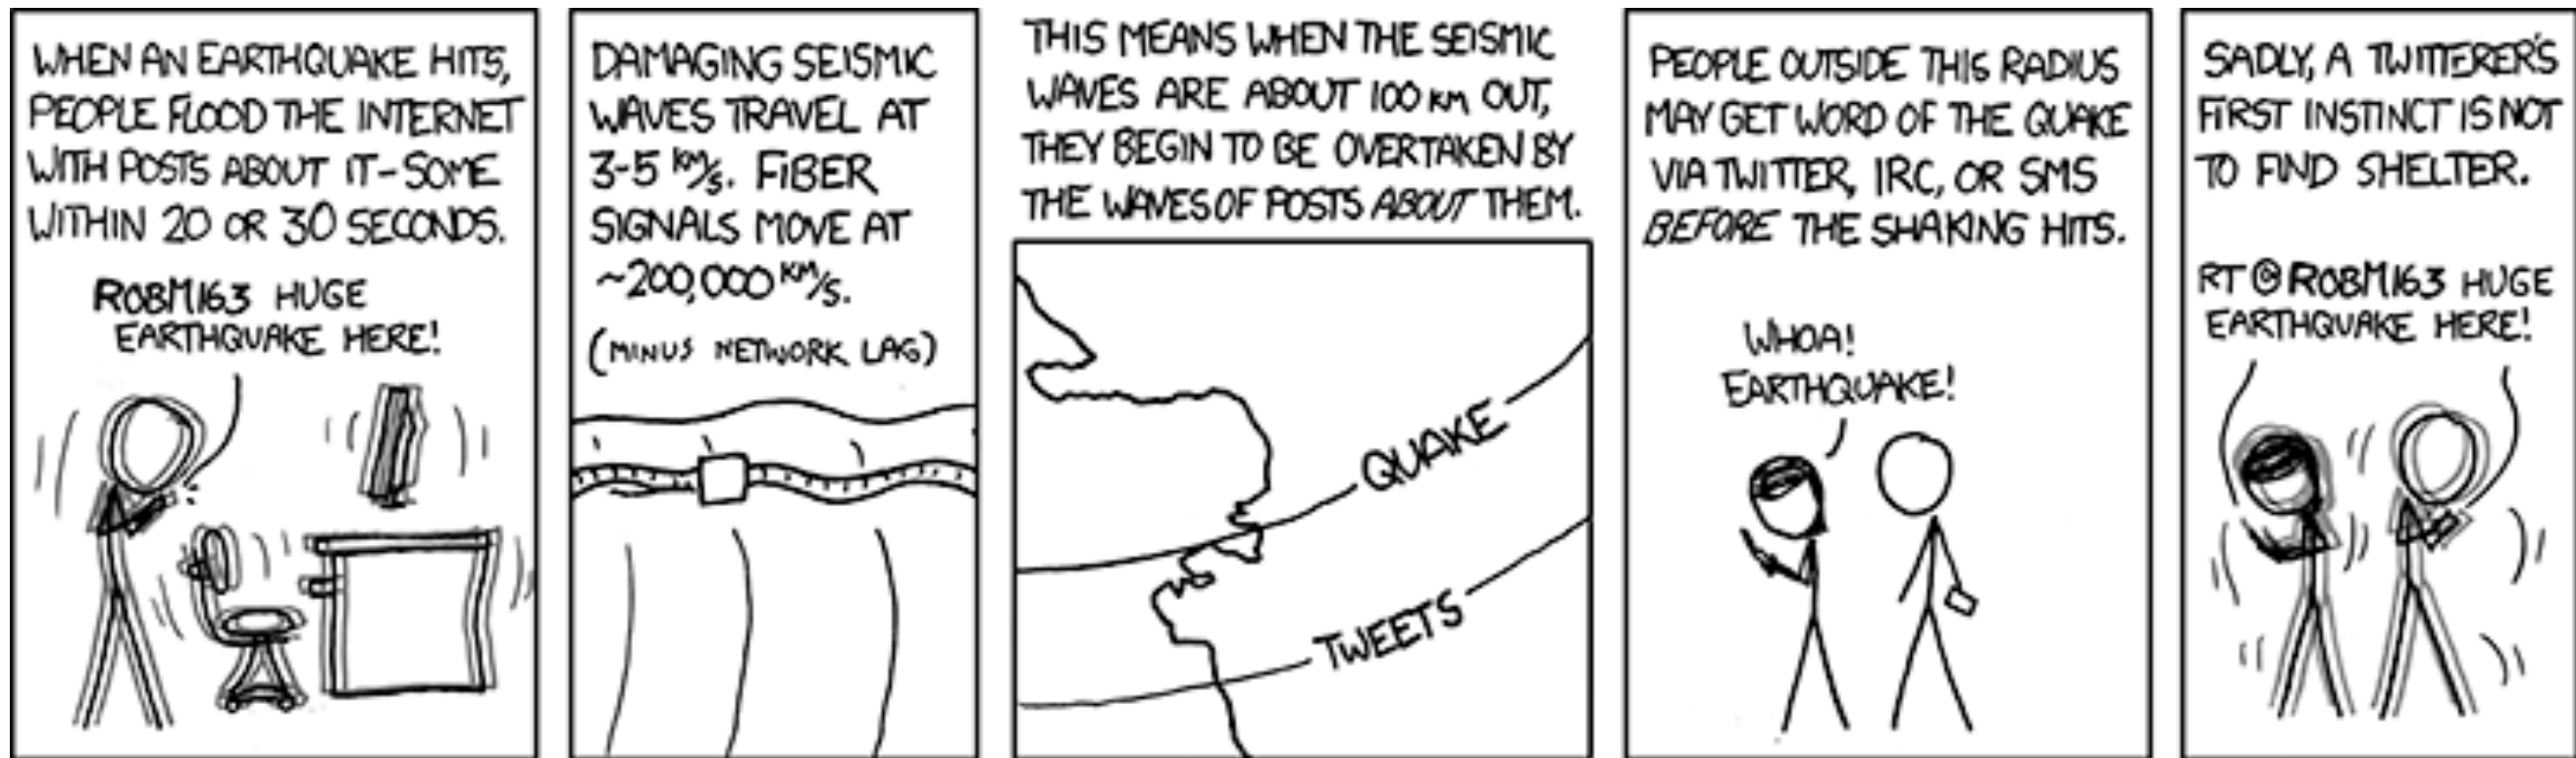


# xkcd.com



# End To End Protocols

COS 460 & 540

# End to End Protocols



This section is about  
**Process to Process**  
communications.

or the **how** applications can talk  
to each other.

# Requirements

- **Guarantee Delivery**
- **Deliver in order**
- **Deliver at most one copy**
- **Any size**
- **Flow control & Synchronization**

# The Network May...

- Drop messages
- Reorder messages
- Deliver duplicate copies
- Limit message size
- Delay messages

# End To End Protocols

The network provides a **best-effort** service



This layer needs to provide  
**high level services**

# End To End Protocols

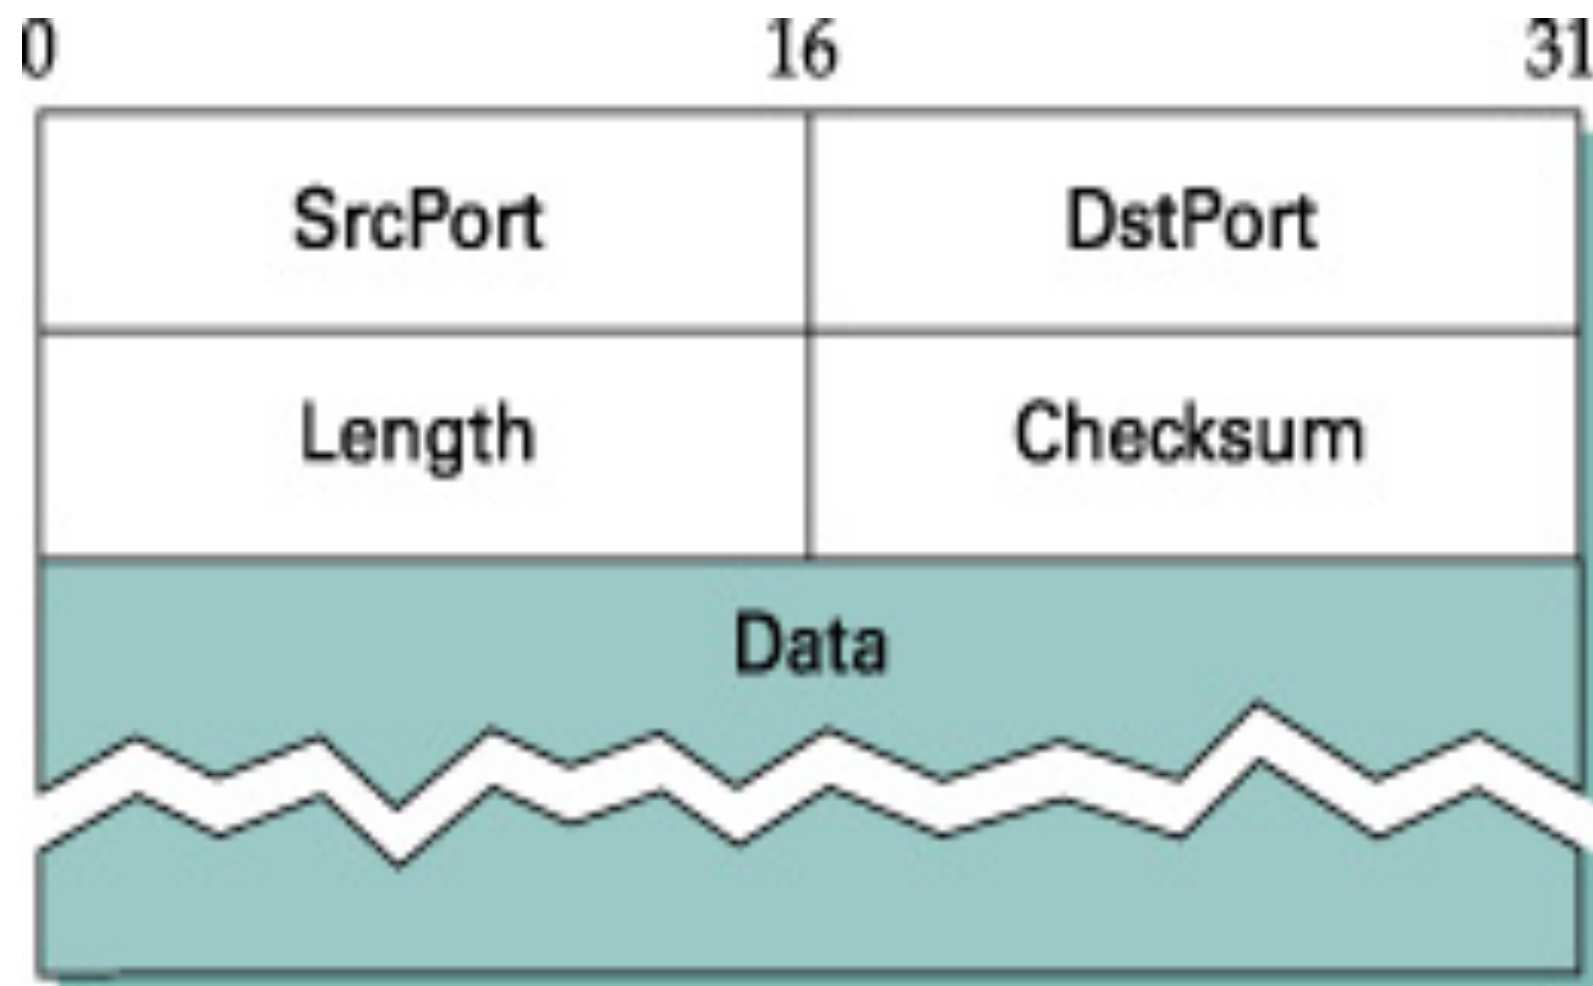
- ★ Simple Demultiplexer (UDP)
- Reliable Stream (TCP)
- Remote Procedure Call (RPC)
- Real-time Applications (RTP)
- Performance

# User Datagram Protocol (UDP)

- Thin layer over Network Layer  
(unreliable, best-effort service)
- Simple demultiplexing; Adds **Process Address** or **port**
- Checksum (packet correctness)



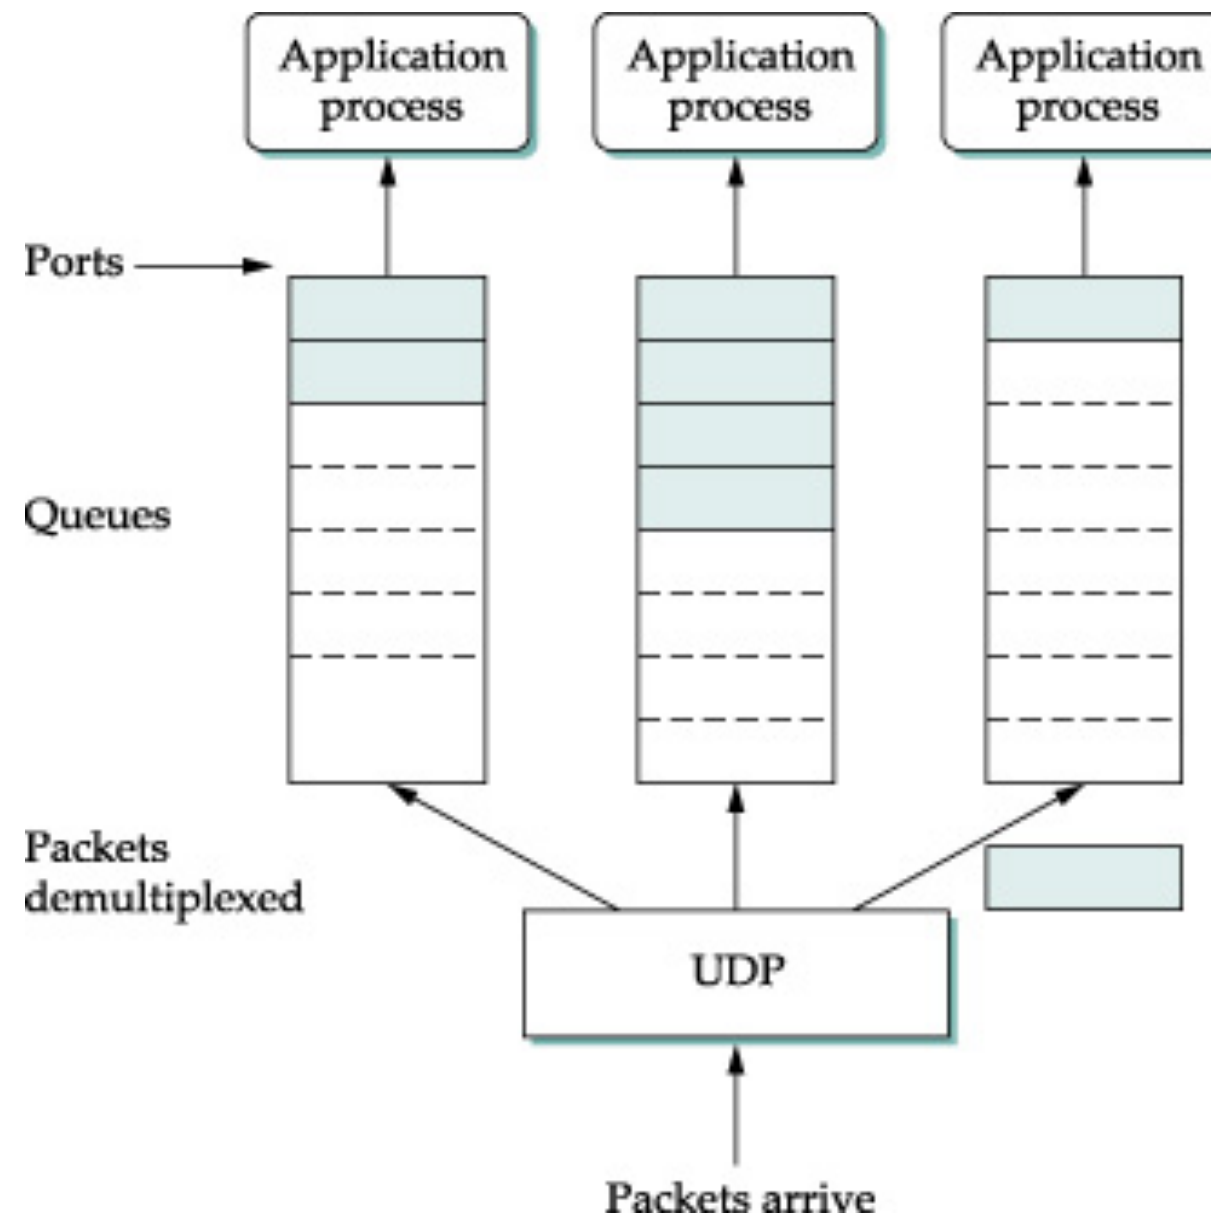
# UDP



# Ports

- 16 bit port number = 64,536 ports
- Unique to a host (IP address)
  - for example: 192.168.1.2:80
- *well-known* ports
  - 22, 25, 80, ...
- <1,024 usually requires administrator

# UDP Demultiplexing



# Socket Pairs

Every datagram has a unique “pairing”

- source ip address
- source port
- destination ip address
- destination port

# End To End Protocols

- ✓ Simple Demultiplexer (UDP)
- ★ Reliable Stream (TCP)
  - Remote Procedure Call (RPC)
  - Real-time Applications (RTP)
  - Performance

# Reliable Byte Stream

Provide a reliable, in-order, at most once delivery mechanism with flow control and synchronization.

... sounds easy right?

# Transmission Control Protocol (TCP)

- Byte-oriented
- Connection-oriented
- Reliable, guaranteed delivery
- In-order delivery
- full-duplex channel

# End-To-End Issues

- Logical connections over the great unknown network
- No “single path,” each packet is routed individually.
- Solve by using Sliding Window Protocol with connection setup and teardown



# End-To-End Issues

- Round Trip Time (RTT) varies widely
  - even over short-lived connections
- Solve by using *adaptive retransmission*

# End-To-End Issues

- Packets are likely to get delivered out of order
- TTL removes packets from the network
- Old packets suddenly show up later
- Solve using Sliding Window w/sequence numbers

# End-To-End Issues

- Bandwidth at endpoints and links between them may be vastly different
- Solve by using a flow control mechanism

# End-To-End Issues

- The Sender may be able to generate large amounts of data
- Intermediate links can become congested
- Solve using congestion control mechanisms

# Segmentation

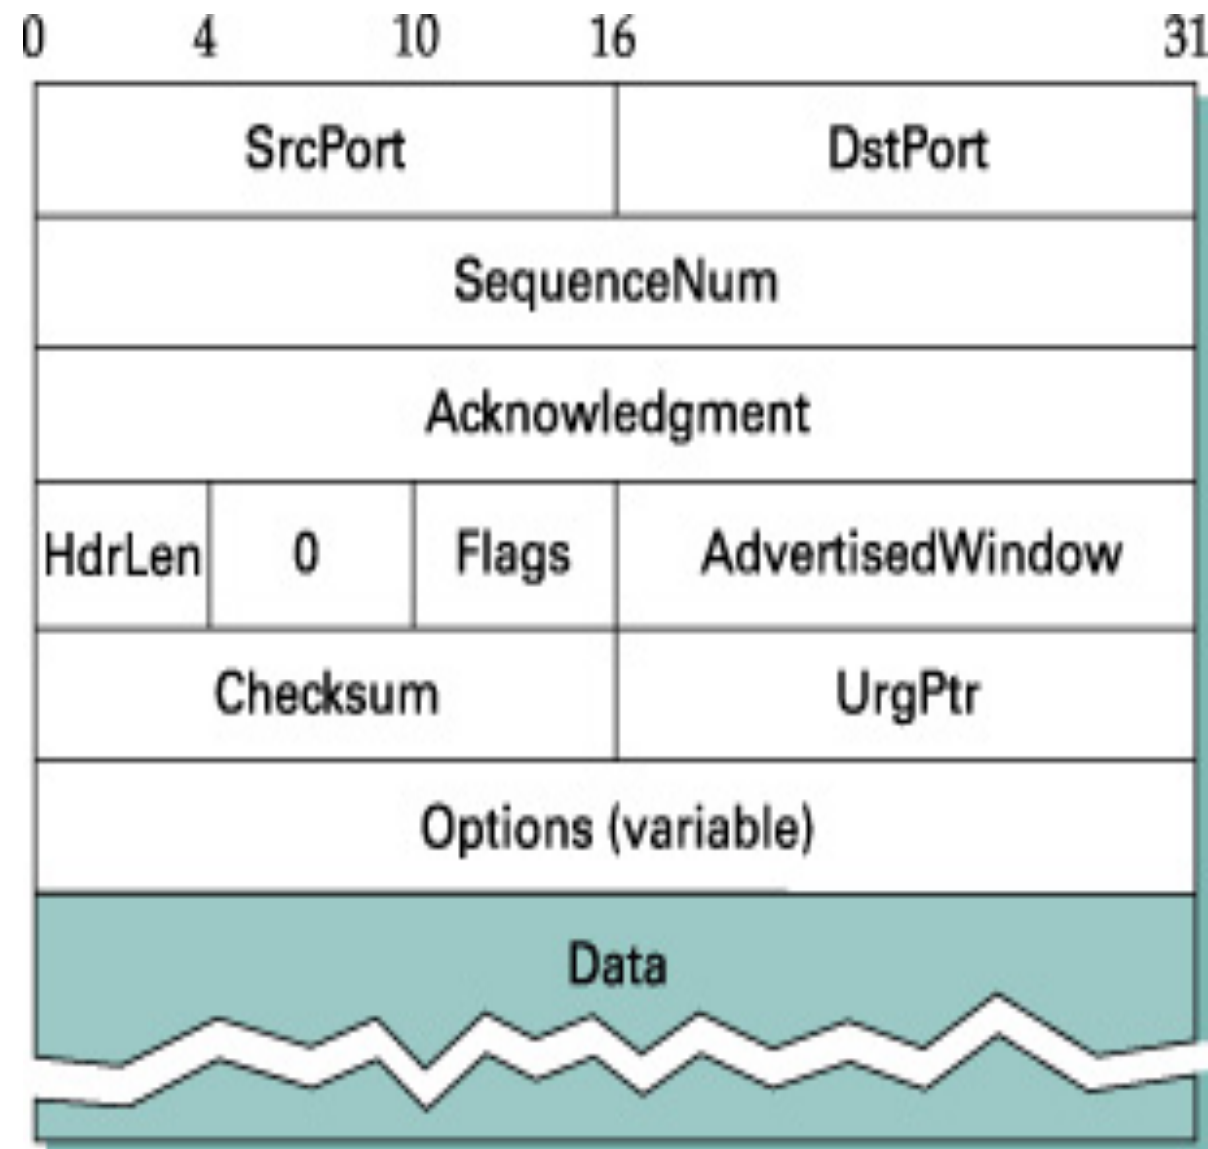
TCP may be a ***byte-stream*** protocol...

...but the network is ***packet based***.


# TCP Segment Format

- TCP is byte-oriented
  - buffers bytes for transmission
  - send when full or flushed
- Send bytes in **segments**
- **Acknowledgments** and **windows**

# TCP Segment Format



# TCP Flow

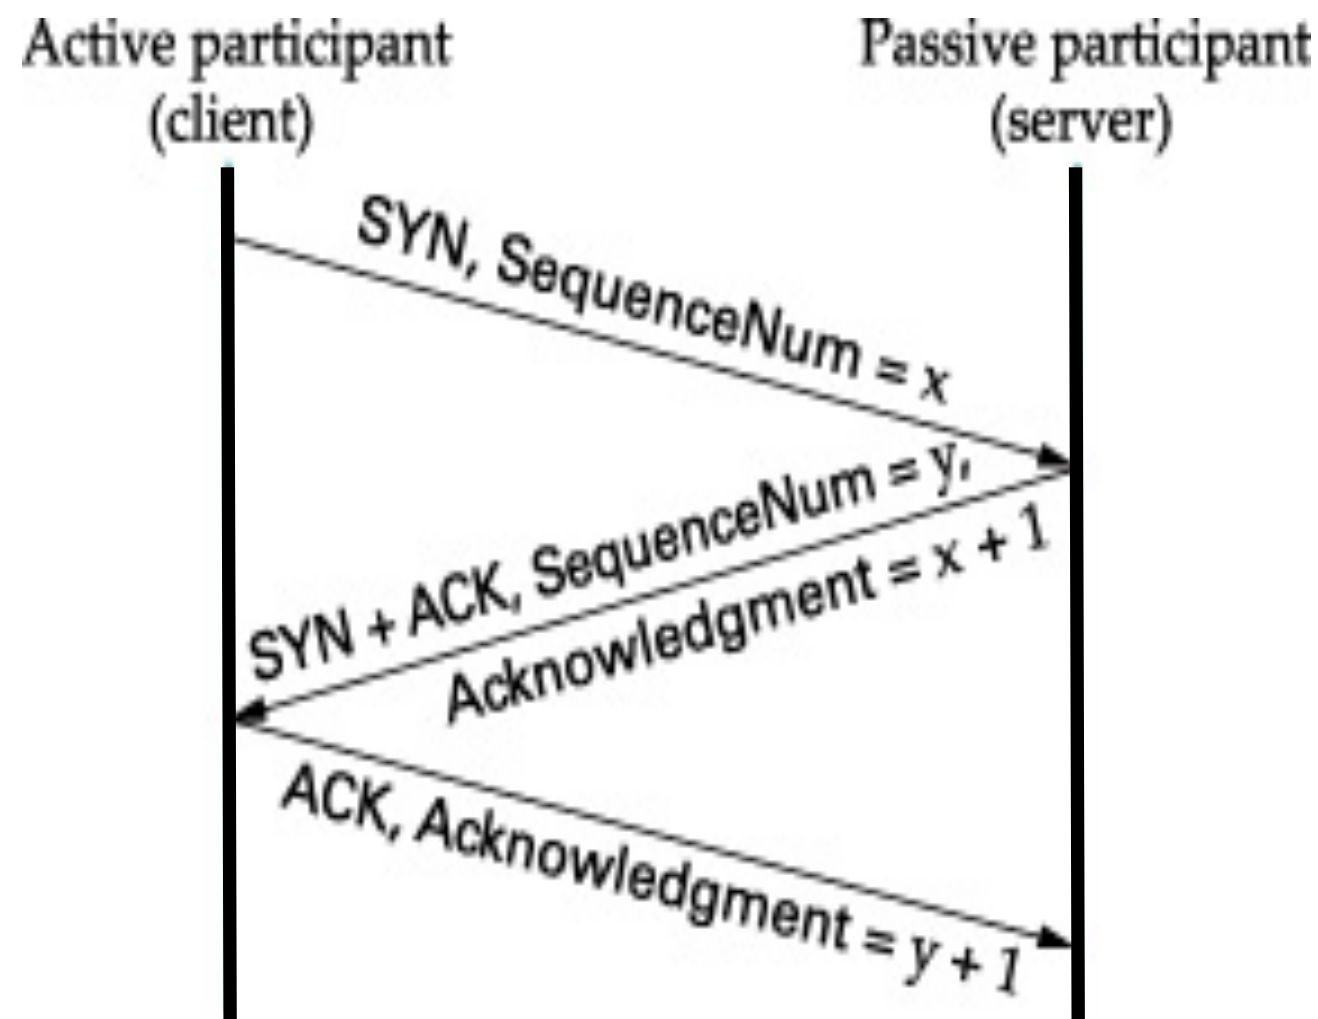




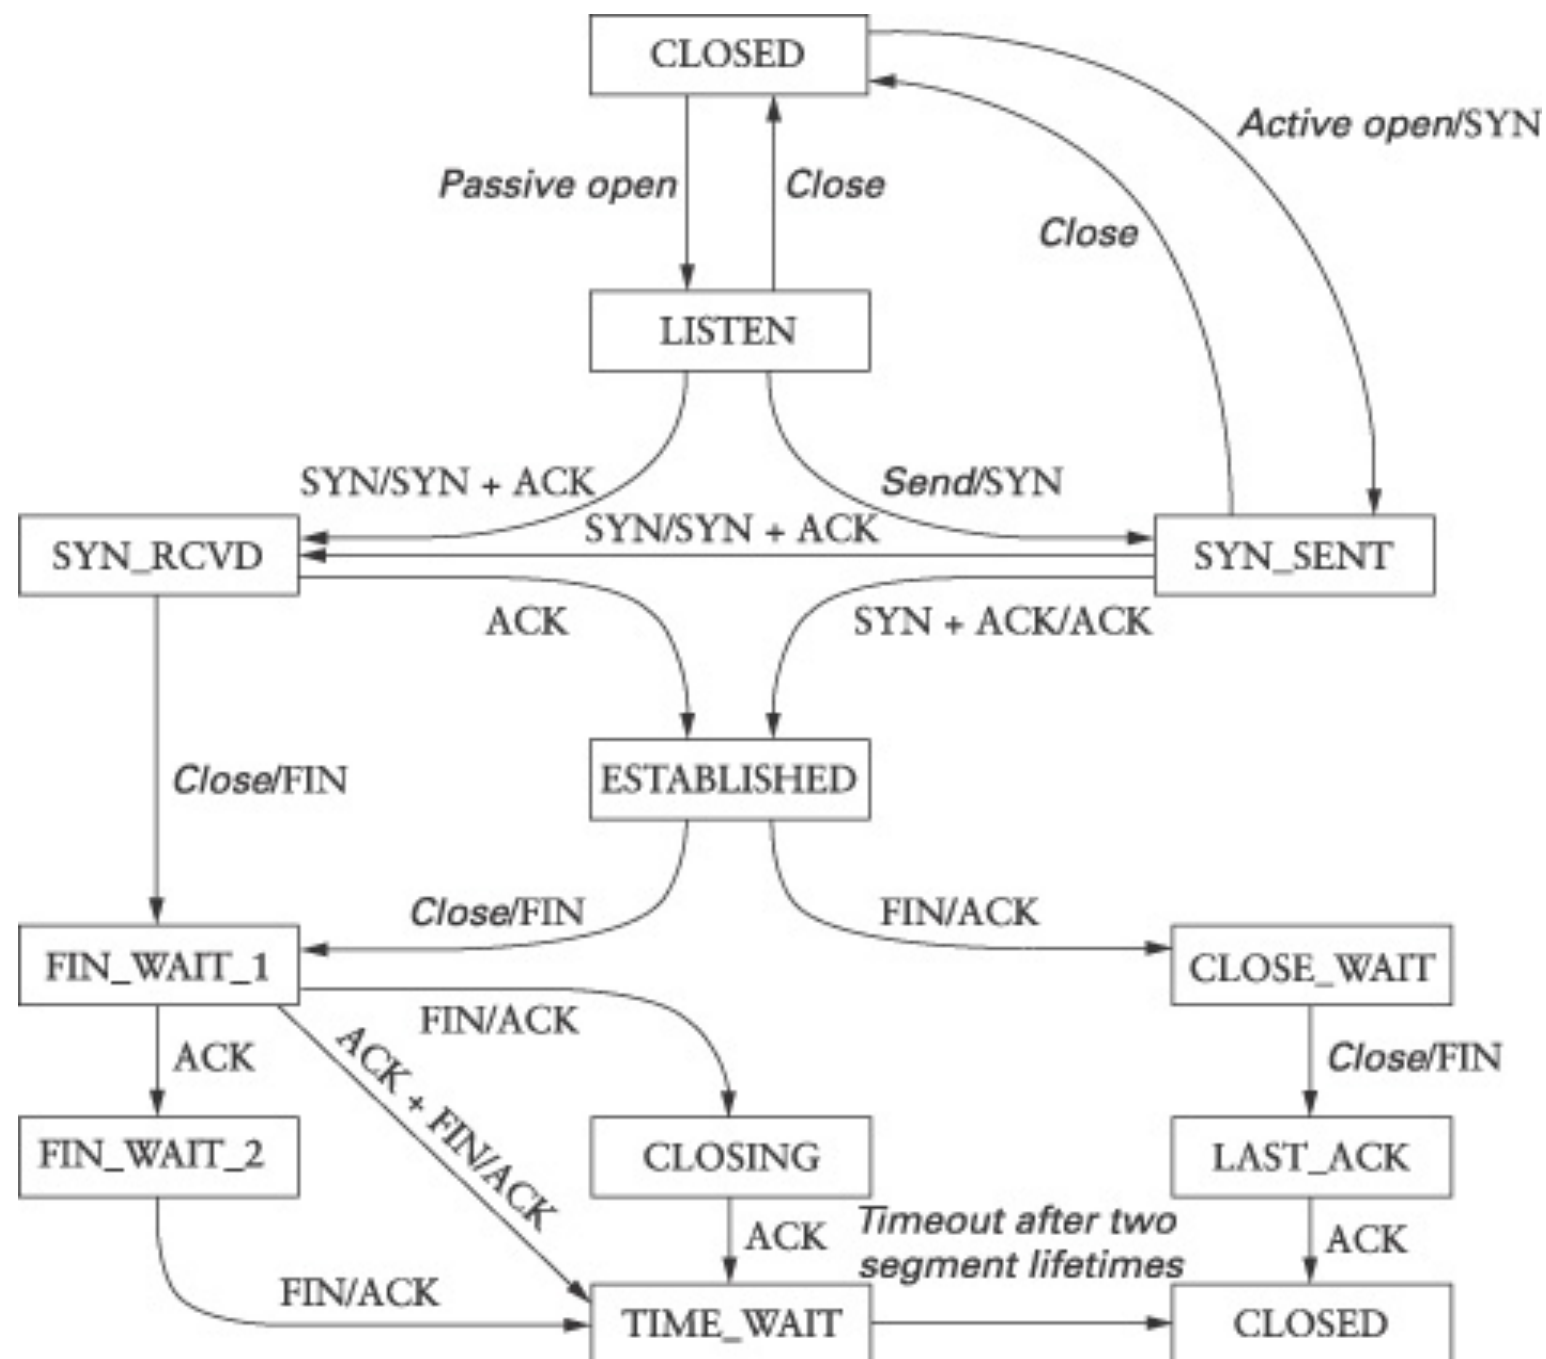
# TCP Segment

- Two way communication
- SequenceNum = first byte in segment
- Acknowledgement = next seq. expected
- AdvertisedWindow = buffer size left

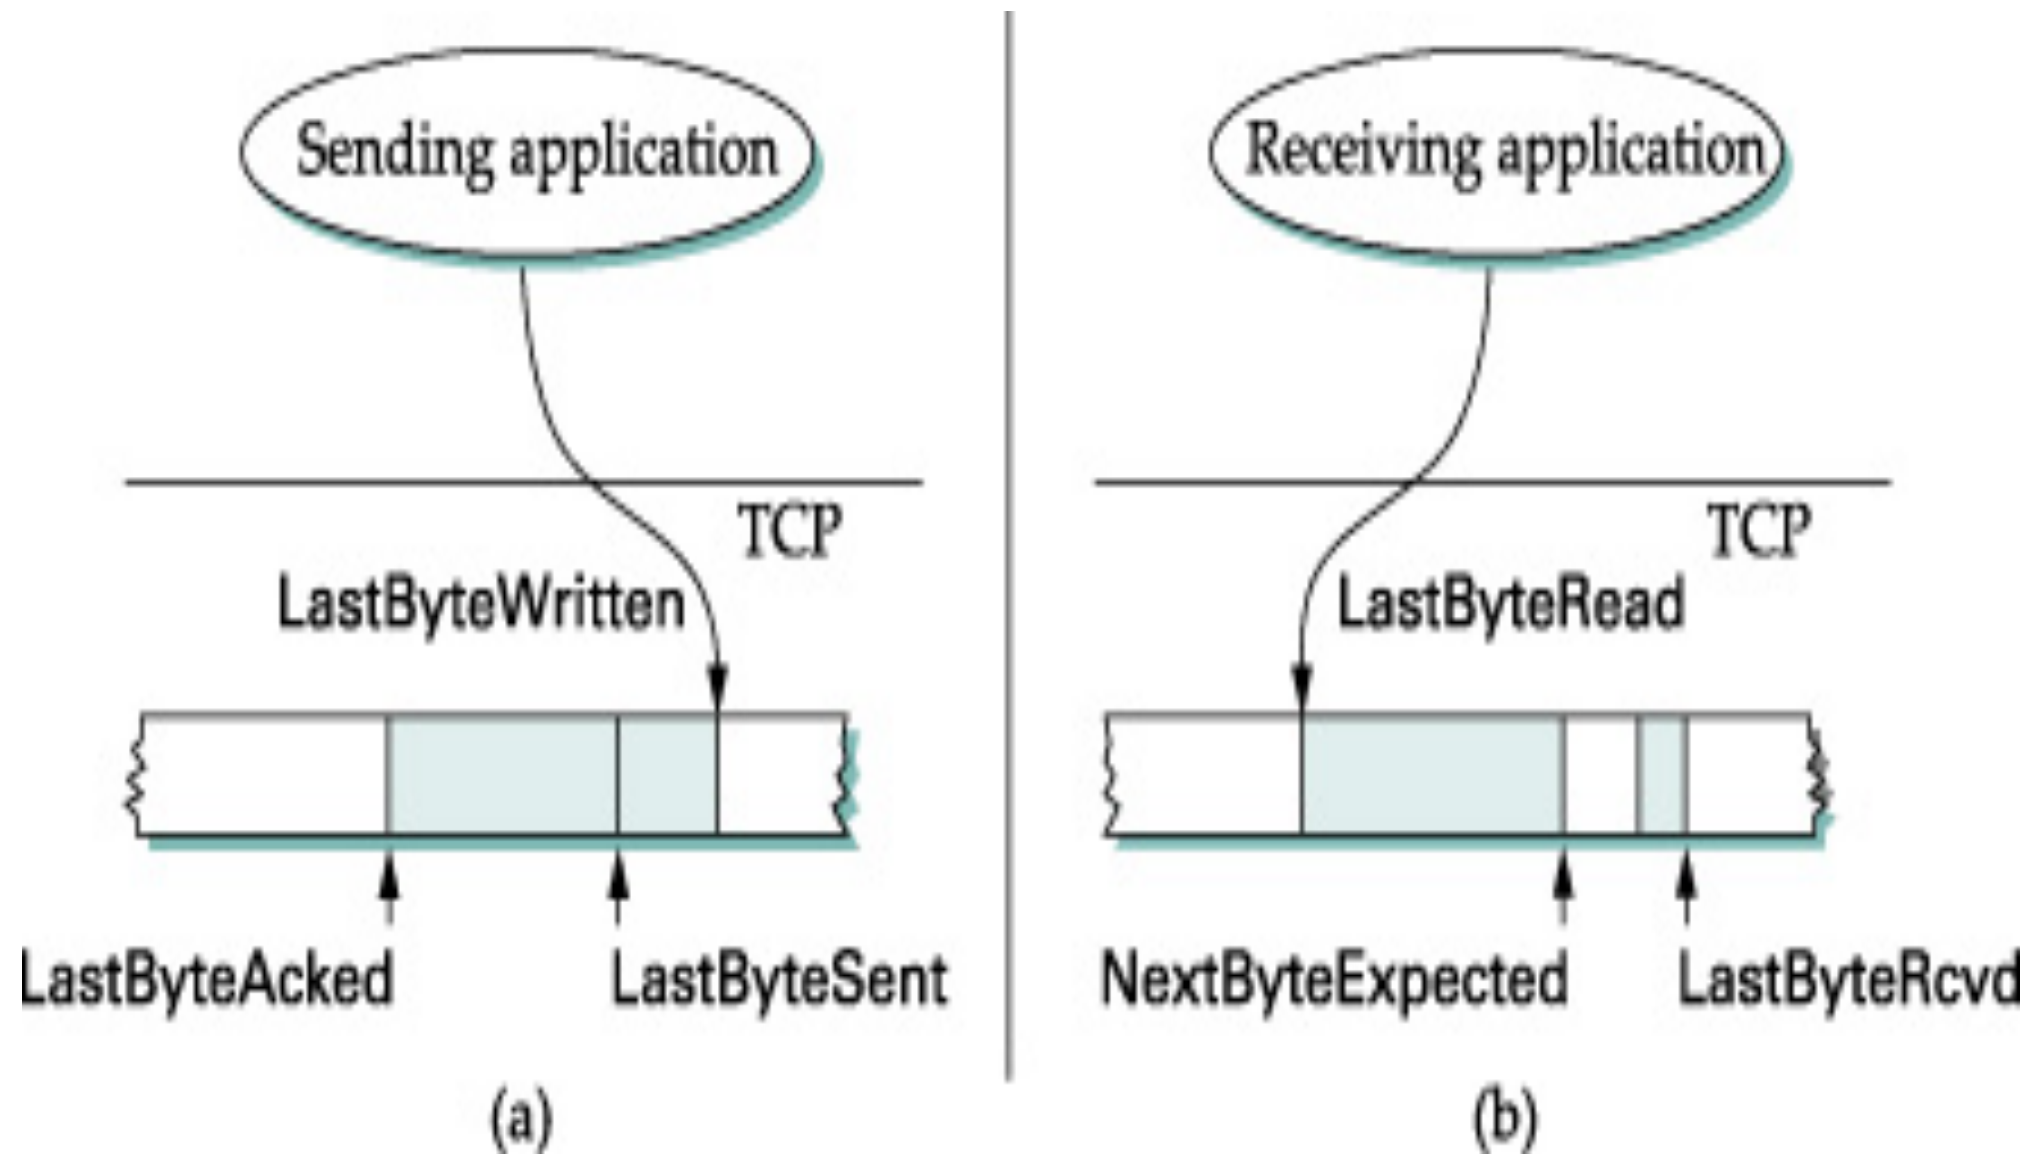
# Connection Establishment



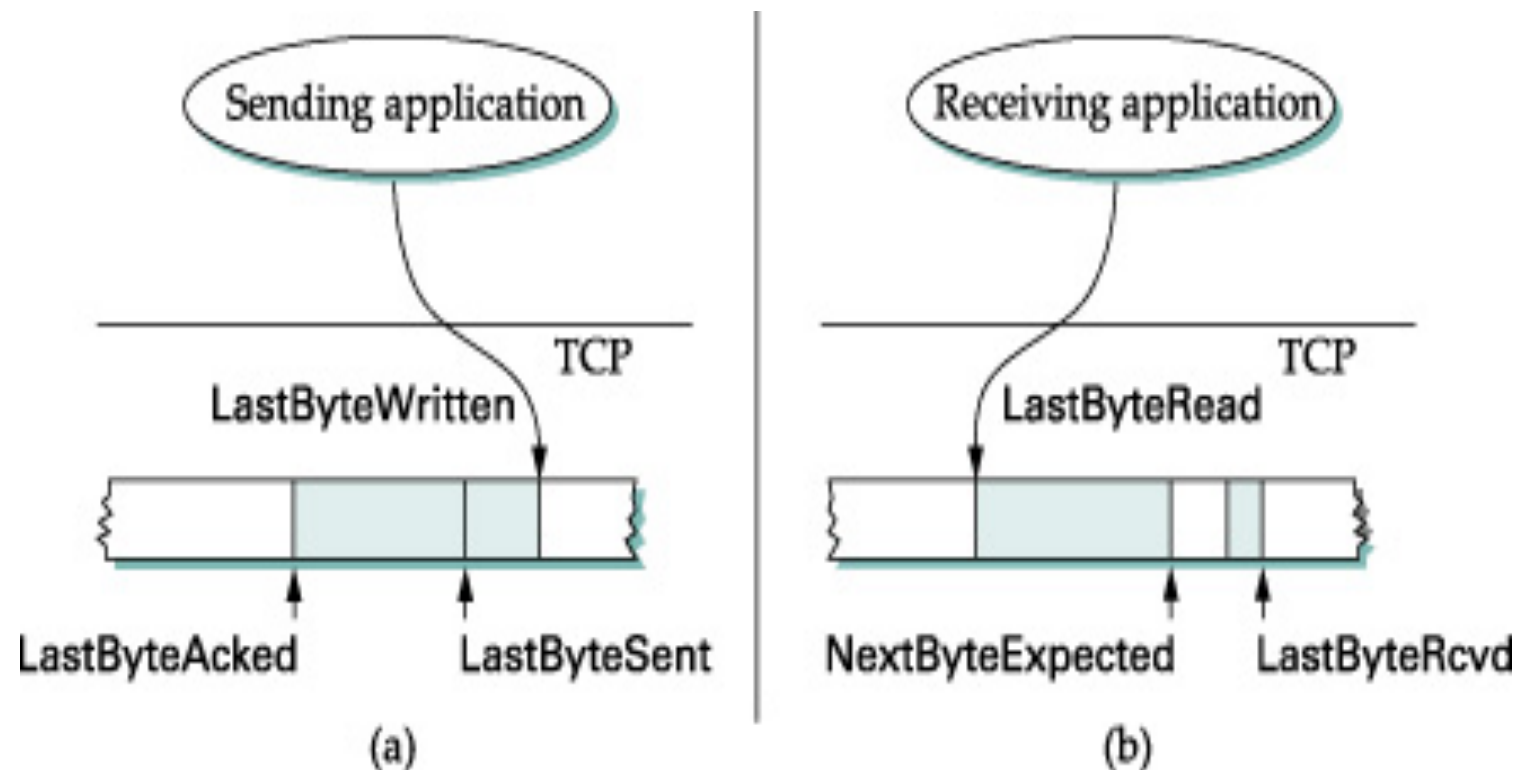
# TCP Lifecycle



# TCP Sliding Window



# TCP Flow Control



Advertised Window

$$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxBuff} =$$
$$\text{MaxB} - ((\text{NextByte} - 1) - \text{LastByteRead})$$

Advertised Window = 0

Sender “pings” with 1 byte segments

# Triggering Transmission

- When segment is full ( $MSS = MTU$ )
- *push* operation
- Timeout expires

# Record Boundaries

How would we send a set of database records?

- <name, address, phone,...>, <...>, ...
- UDP = 1 record per datagram

# Record Boundaries

In TCP we could....

- Encode special characters, recall ETX?
- Encode the dataset, e.g. XML
- Urgent data flag
- push operation



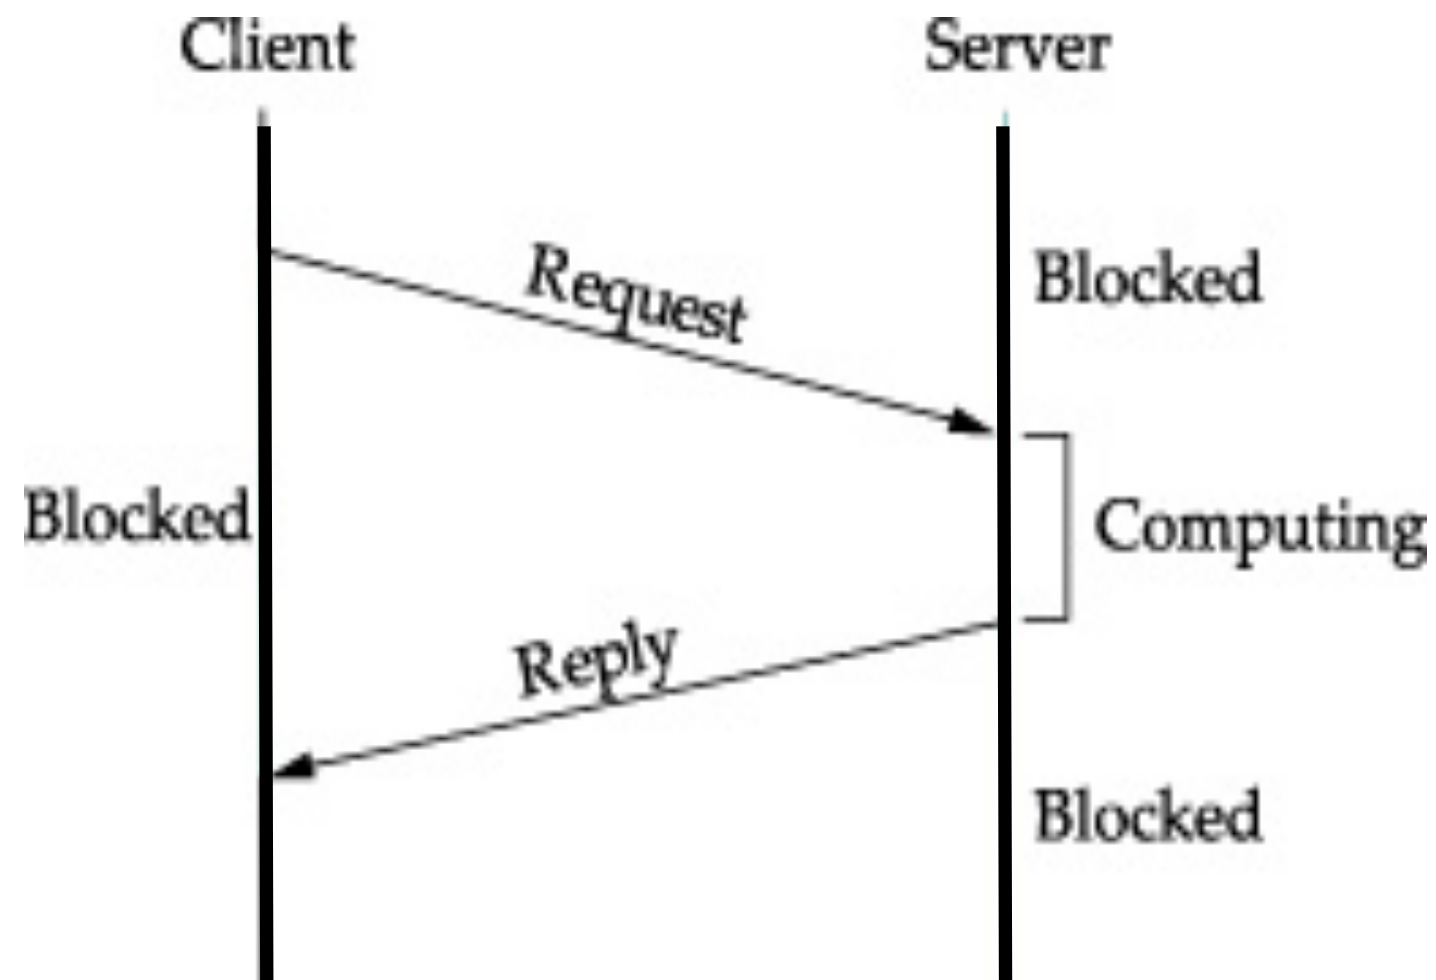
# End To End Protocols

- ✓ Simple Demultiplexer (UDP)
- ✓ Reliable Stream (TCP)
- ★ Remote Procedure Call (RPC)
  - Real-time Applications (RTP)
  - Performance

# Remote Procedure Call

- Request/Reply Protocol
- Different from UDP?
  - Reliable
  - Sequencing

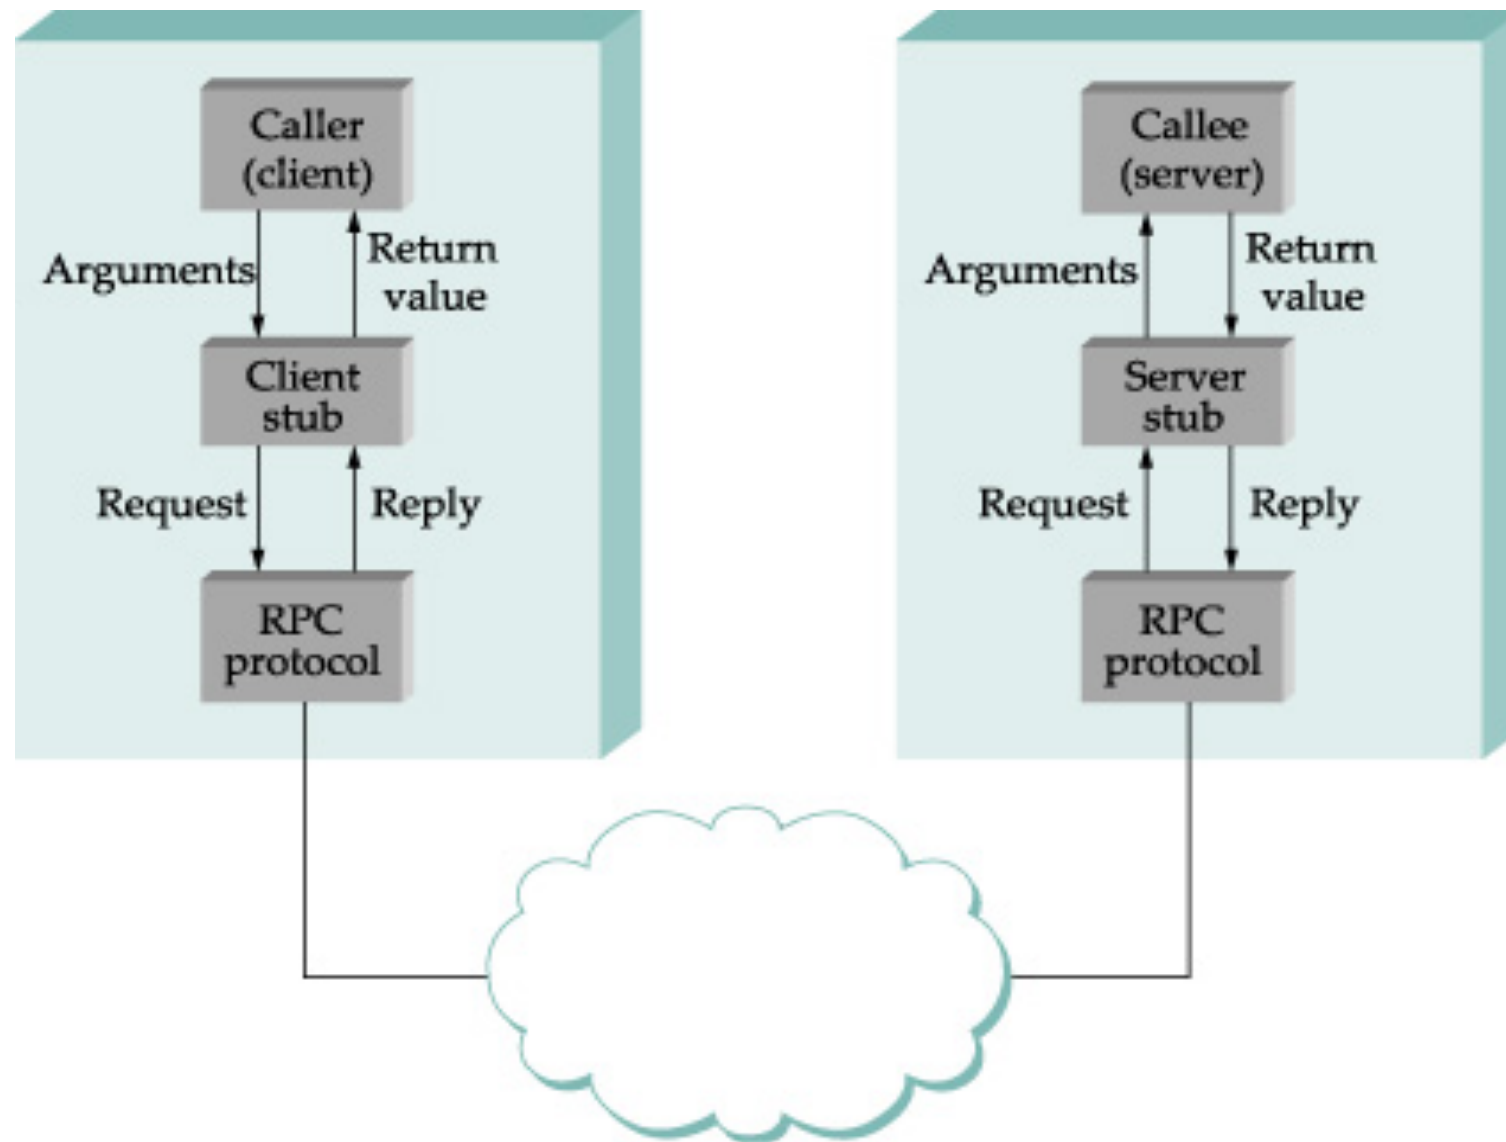
# Request Reply



# RPC Stack

- Client-Side
  - Generates “stubs”
  - Developer calls stubs
- Server-Side
  - Generates “server”
  - Server calls developers routines

# RPC Stack



# Serializing Data

- byte, word, character
- string
- array
- tree
- ... more complex data structures

# RPC Questions

- RPC over UDP or over TCP or over IP?
- Textual data representation or binary
  - e.g. XML vs XDR

# Other RPC-like Protocols

- Common Object Request Broker (CORBA)
- Distributed Component Object Module (DCOM)
- Remote Method Invocation (RMI)
- Simple Object Access Protocol (SOAP)



# End To End Protocols

- ✓ Simple Demultiplexer (UDP)
- ✓ Reliable Stream (TCP)
- ✓ Remote Procedure Call (RPC)
- ★ Real-time Applications (RTP)
- Performance

# Real-Time Applications

- Encodings... (mp3, aac, etc..)
- Timing of data at receiver
  - Synchronization
- Packet loss
- Frame boundaries

# Real-Time Applications

- TCP
  - Has it's own reliable transmission
  - Does not “know” the data
- UDP
  - Does very little, works well

# End To End Protocols

- ✓ Simple Demultiplexer (UDP)
- ✓ Reliable Stream (TCP)
- ✓ Remote Procedure Call (RPC)
- ✓ Real-time Applications (RTP)
- ★ Performance

# End To End Protocols

- ✓ Simple Demultiplexer (UDP)
- ✓ Reliable Stream (TCP)
- ✓ Remote Procedure Call (RPC)
- ✓ Real-time Applications (RTP)
- ✓ Performance

# End

End To End Protocols