

1

Foundations

COS 460 & COS 540

2

Foundations



- Applications
- Requirements
- Architecture
- Performance
- Writing Code

3

Applications

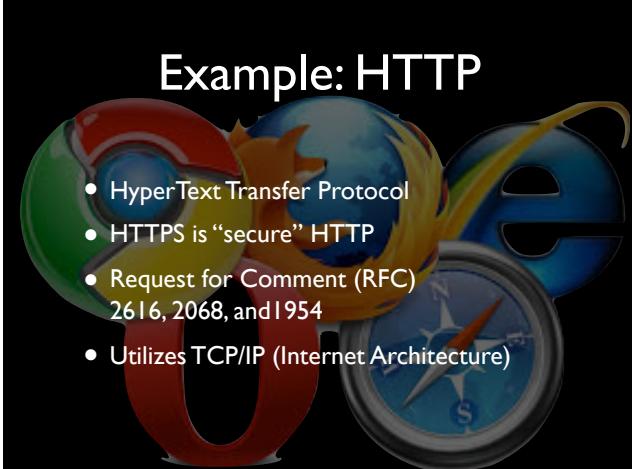


...and these are just the web “social media” ones

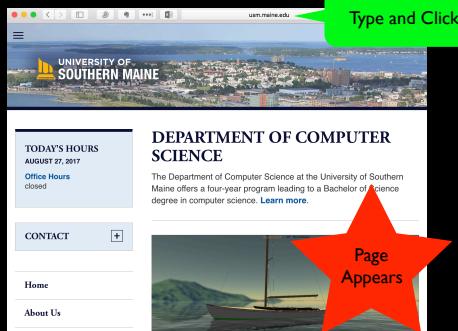
4

Example: HTTP

- HyperText Transfer Protocol
- HTTPS is “secure” HTTP
- Request for Comment (RFC) 2616, 2068, and 1954
- Utilizes TCP/IP (Internet Architecture)

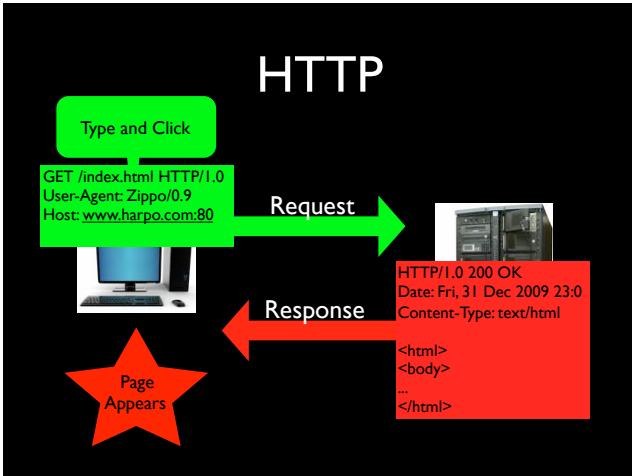


5



How does this work?

6



HTTP,TCP



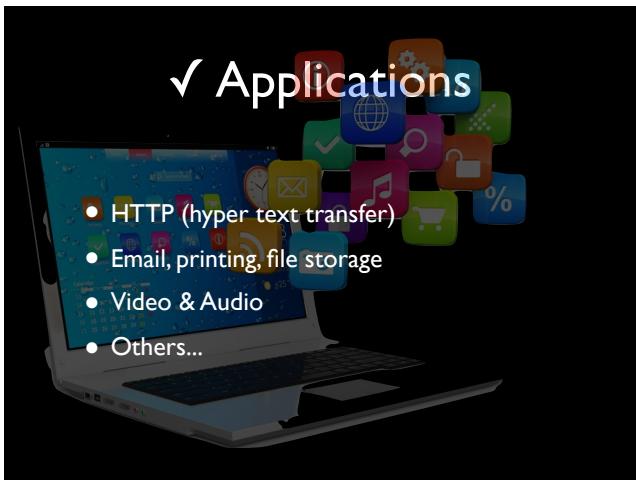
HTTP

- HTTP is essentially **file transfer**
- Every bit is important
- Similar to Email, Printing, File Servers...

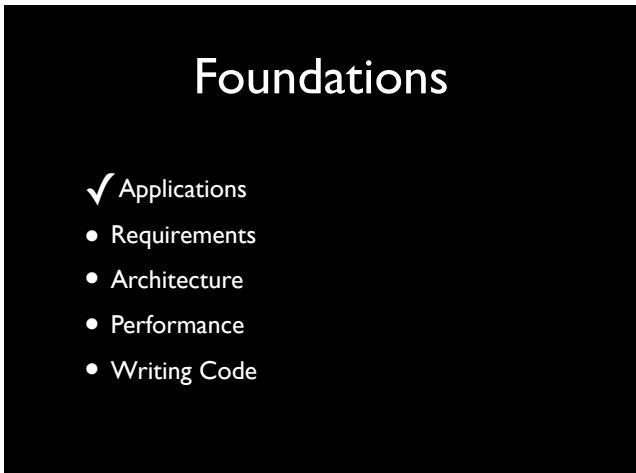
Audio & Video

Every bit is not
important but
their **timing** is





10



11

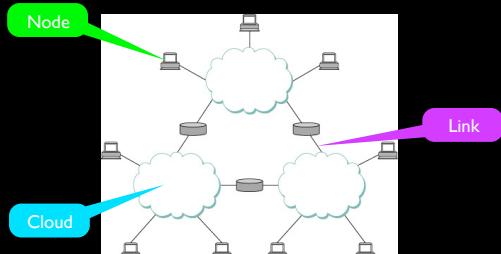


12

Who's Requirements?

- Application Programmer
- Network Provider (ISP)
- Network Designer

Connectivity

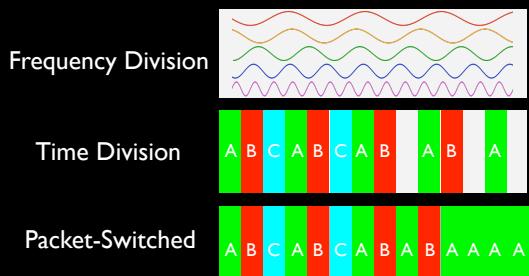


Cost-Effective

How do we get as **much information** as possible, from **multiple** sources to multiple sources across the “wire”?

...one bit at a time is slow.

Multiplexed



16

Imagine A has 2x B's data needs and B has 2x C's
FDM – Unused allocation for D and E as well as B and C not using all
TDM – Imagine A has more data, unused time slots
PSwitched – Quality of Service & Congestion problems

Common Services

Define **useful channels** that understand the **application needs** and the **networks ability**.

17

Common Services

- File transfer vs. Streaming
- Reliable vs. Unreliable
- Bandwidth vs. Latency
- Large vs. Small

18

19

✓ Requirements

- **Connectivity**
Links, Nodes, Clouds
- **Cost-Effective**
Packet Switched
- **Support (common) Services**
File transfer, streaming, bandwidth...

20

Foundations

- ✓ Applications
- ✓ Requirements
- Architecture
- Performance
- Writing Code

21

Code

It Takes Two to Tango

- Server (passive)
- Client (active)



22

Server

- **Binds** to a “port”
- Passively **listens** for connection request
 - **Accepts** connections
 - reads and writes data
 - **Closes** connection when done

23

Client

- Actively **connects** to server's port
 - reads and writes on connection
 - **Closes** connection when done

24

Java Server

```
// passive bind and listen  
ServerSocket srv = new ServerSocket(port);  
Socket s = srv.accept();  
    s.getInputStream() & read()  
    s.getOutputStream() & write()  
s.close();
```

Java Client

```
// active connect  
Socket s = new Socket("server", port);  
    s.getOutputStream() & write()  
    s.getInputStream() & read()  
s.close()
```

✓ Code

- Client (active) - Server (passive)
- Peer to Peer networks?
 - Who's the server?
- Distributed networks
 - Multiple servers and clients

Foundations

- ✓ Applications
- ✓ Requirements
- ✓ Architecture
- ✓ Performance
- ✓ Writing Code

28

Architecture

How do we build a Network so it is...

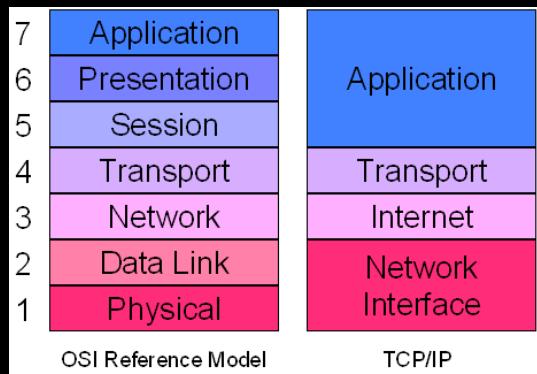
- understandable
- programmable
- flexible
- useful

29

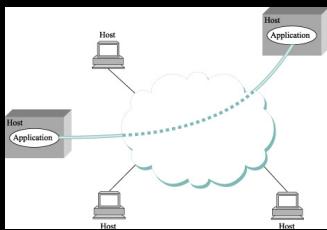


Layers make everything better

30

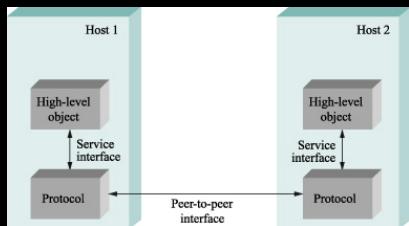


31



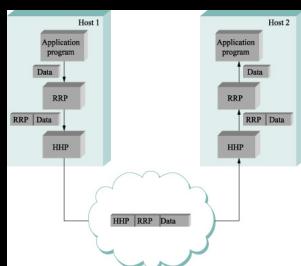
Application View

32



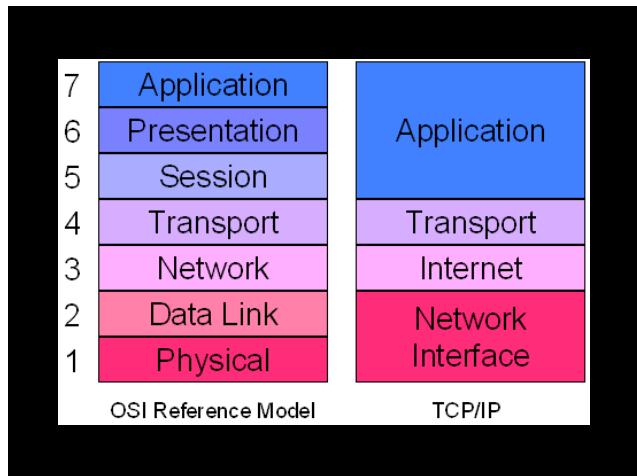
Service/Peer Interfaces

33



Encapsulation

34



35

✓ Architecture

Network Layers

- Subdivide the problem
- Address one at a time
- Do a good job at each layer

36

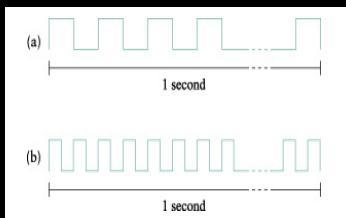
Foundations

- ✓ Applications
- ✓ Requirements
- ✓ Architecture
 - Performance
 - Writing Code

Performance

- Bandwidth
- Latency
- Delay X Bandwidth
- High-Speed Networks

Bandwidth



Bits transmitted per unit of time

Bandwidth

- 10 Mbps
 - 10 million bits per second
 - $0.1\mu\text{s}$ for each bit
- 20 Mbps; $0.5\mu\text{s}$ for each bit
- 100 Mbps; $0.01\mu\text{s}$ for each bit

40

Latency

How long does it take to get there?



41

Latency

- = Propagation + Transmit + Q-Delay
- Propagation Delay (distance / c)
- Transmit Time (size / bandwidth)
- Queueing Delays

1-byte to/from server, 1Mbps
 $\text{Transmit} = 1\mu\text{s}/\text{bit} = 8\mu\text{s}$
 1-byte to/from server, 100Mbps
 $\text{Transmit} = 0.01\mu\text{s}/\text{bit} = 0.08\mu\text{s}$
25MB file @ 10Mbps
 $25 \times 10^6 \times 8\text{bits} / 10 \times 10^6 \text{ Mbps} = 20\text{sec}$

42

Latency

- Speed of Light (c)
 - copper = $2.3 \times 10^8 \text{ m/s}$
 - vacuum = $3.0 \times 10^8 \text{ m/s}$
 - fiber = $2.0 \times 10^8 \text{ m/s}$

43

Latency

- One Way Latency
- Round Trip Time (RTT)

44

Bandwidth vs Latency

For...

- HTTP?
- File Transfer?
- Audio?
- Video?

45

Bandwidth vs Latency

Consider a 1-byte message and a 1-byte response over a 10Mbps link with a 10ms RTT.

...Does **bandwidth** or **latency** dominate the transmission?
...When will bandwidth dominate?

$0.8\text{us} + 10\text{ms} = 10.0008\text{ms}$ —
Latency dominates
Even at 1Gbps — $0.008\text{us} + 10\text{ms} = 10.00008$ Latency dominates

@10Mbps; 10,000bits = 1ms;
100,000bits = 10ms; 12,500 bytes or ~12kB

Delay x Bandwidth

Amount of data “in-flight”

$$\begin{aligned}
 & 10 \text{ Mbps} \times 50 \text{ mS} \\
 & (10 \times 10^6 \text{ bits/sec} \times (50 \times 10^{-3}) \text{ sec}) \\
 & 500 \times 10^3 = 500\text{Kb} \\
 & 62.5\text{KB}
 \end{aligned}$$

High-Speed Networks

- Very large **bandwidth**
- Latency does **not** change
- Delay X Bandwidth goes **up**
- More data in-flight during RTT
- Latency begins to dominate transfer times.

✓ Performance

- Bandwidth
- Latency
- Delay x Bandwidth
- High-Speed networks

Foundations

- ✓ Applications
- ✓ Requirements
- ✓ Architecture
- ✓ Performance
- Writing Code

Foundations

End

```
# Exercise 1.4
1.5MB file,
80ms RTT,
1KB packet size,
2 RTT handshake at start

size = (1.5 * 2^20) * 8 => 12,582,912
rtt = 0.080
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.16
```

```

# Exercise 1.4
1.5MB file, 80ms RTT, 1KB packet size, 2 RTT handshake
at start

size = (1.5 * 2^20) * 8 => 12,582,912
rtt = 0.080
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.16

a) 10Mbps bandwidth, send continuous

bandwidth = 10 * 10^6 => 10,000,000

propagation_delay = rtt / 2 => 0.04
transmit_time = size / bandwidth => 1.2583

handshake + propagation_delay + transmit_time
=> 1.4583

```

52

```

# Exercise 1.4
1.5MB file, 80ms RTT, 1KB packet size, 2 RTT handshake
at start

size = (1.5 * 2^20) * 8 => 12,582,912
rtt = 0.080
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.16

b) 10Mbps bandwidth, wait 1 RTT after each packet

bandwidth = 10 * 10^6 => 10,000,000

propagation_delay = rtt / 2 => 0.04
transmit_time = size / bandwidth => 1.2583

packet_count = size / packet_size => 1,536
t_time = transmit_time + ((packet_count - 1) * rtt)
=> 124.0583

handshake + propagation_delay + t_time => 124.2583

```

53

```

# Exercise 1.4
1.5MB file, 80ms RTT, 1KB packet size, 2 RTT handshake
at start

size = (1.5 * 2^20) * 8 => 12,582,912
rtt = 0.080
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.16

b) Alternate version

bandwidth = 10 * 10^6 => 10,000,000
propagation_delay = rtt / 2 => 0.04

packet_time = (packet_size / bandwidth) + rtt
=> 0.0808
t_time = (packet_count - 1) * packet_time
=> 124.0575

handshake + propagation_delay + t_time => 124.2575

```

54

```
# Exercise 1.4  
1.5MB file, 80ms RTT, 1KB packet size, 2 RTT handshake  
at start
```

```
size = (1.5 * 2^20) * 8 => 12,582,912  
rtt = 0.080  
packet_size = 2^10 * 8 => 8,192  
handshake = 2 * rtt => 0.16
```

c) infinite transmit, 20 packets per RTT

```
packet_count = size / packet_size => 1,536  
packet_chunks = ceil(packet_count / 20) => 77  
  
handshake + (packet_chunks * rtt) => 6.32
```

Don't need to wait for response on last packet to
calculate arrival times.

```
handshake + (packet_chunks * rtt)  
- propagation_delay => 6.28
```

55

```
# Exercise 1.4  
1.5MB file, 80ms RTT, 1KB packet size, 2 RTT handshake  
at start
```

```
size = (1.5 * 2^20) * 8 => 12,582,912  
rtt = 0.080  
packet_size = 2^10 * 8 => 8,192  
handshake = 2 * rtt => 0.16
```

d) 0 transmit time, 1 packet first RTT, 2 second RTT, 4
third...

```
packet_count = size / packet_size => 1,536  
packet_chunks = ceil(log(packet_count, 2)) => 11
```

Don't need to wait for response on last packet to
calculate arrival times.

```
handshake + (packet_chunks * rtt)  
- propagation_delay => 1
```

56

```
# Exercise 1.3  
1,000KB file,  
50ms RTT,  
1KB packet size,  
2 RTT handshake at start
```

```
size = (1,000 * 2^10) * 8 => 8,192,000  
rtt = 0.050 => 0.05  
packet_size = 2^10 * 8 => 8,192  
handshake = 2 * rtt => 0.1
```

57

58

```
# Exercise 1.3
1,000KB file, RTT 50ms, 2 RTT handshake at start

size = (1,000 * 2^10) * 8 => 8,192,000
rtt = 0.050 => 0.05
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.1

a) 1.5Mbps, continuous

bandwidth = 1.5 * 10^6 => 1,500,000
propogation_delay = rtt / 2 => 0.025
transmit_time = size / bandwidth => 5.4613

handshake + propogation_delay + transmit_time
=> 5.5863
```

59

```
# Exercise 1.3
1,000KB file, RTT 50ms, 2 RTT handshake at start

size = (1,000 * 2^10) * 8 => 8,192,000
rtt = 0.050 => 0.05
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.1

b) 1.5Mbps, one RTT after each packet

bandwidth = 1.5 * 10^6 => 1,500,000
propogation_delay = rtt / 2 => 0.025
transmit_time = size / bandwidth => 5.4613

packet_count = size / packet_size => 1,000
t_time = transmit_time + ((packet_count - 1) * rtt)
=> 55.4113

handshake + propogation_delay + t_time => 55.5363
```

60

```
# Exercise 1.3
1,000KB file, RTT 50ms, 2 RTT handshake at start

size = (1,000 * 2^10) * 8 => 8,192,000
rtt = 0.050 => 0.05
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.1

c) infinite bandwidth, 20 packets per RTT

packet_count = size / packet_size => 1,000
packet_chunks = ceil(packet_count / 20) => 50

handshake + (packet_chunks * rtt) => 2.6
handshake + (packet_chunks * rtt)
- propogation_delay => 2.575
```

```
# Exercise 1.3
1,000KB file, RTT 50ms, 2 RTT handshake at start

size = (1,000 * 2^10) * 8 => 8,192,000
rtt = 0.050 => 0.05
packet_size = 2^10 * 8 => 8,192
handshake = 2 * rtt => 0.1

d) exponential packet counts, 1, 2, 4, 8, ...

packet_count = size / packet_size => 1,000
packet_chunks = ceil(log(packet_count, 2)) => 10

Don't need to wait for response on last packet to
calculate arrival times.

handshake + (packet_chunks * rtt)
- propagation_delay => 0.575
```